

Homework 2: Actor-Critic, Value Function Approximation, and TRPO

Submission Guidelines: Please compress all your write-ups (photos/scanned copies are acceptable; please make sure that the electronic files are of good quality and reader-friendly) and source code into one .zip file and submit the compressed file via E3.

Problem 1 (Surrogate Function in TRPO)

(15 points)

In this problem, we will prove some key properties of the surrogate function used in TRPO. Recall from Page 27 of the slides of Lecture 14: Assume that both the state space and the action space are finite. The surrogate function $L_{\pi_{\theta_1}}(\pi_\theta)$ is defined as

$$L_{\pi_{\theta_1}}(\pi_\theta) := \eta(\pi_{\theta_1}) + \sum_s d_\mu^{\pi_{\theta_1}}(s) \sum_a \pi_\theta(a|s) A^{\pi_{\theta_1}}(s, a). \quad (1)$$

Show that $L_{\pi_{\theta_1}}(\pi_\theta)$ satisfies the two properties: (i) $L_{\pi_{\theta_1}}(\pi_{\theta_1}) = \eta(\pi_{\theta_1})$ and (ii) $\nabla_\theta L_{\pi_{\theta_1}}(\pi_\theta)|_{\theta=\theta_1} = \nabla_\theta \eta(\pi_\theta)|_{\theta=\theta_1}$.

Problem 2 (Solving TRPO Under Approximation Using Duality)

(15+15=30 points)

Recall from Page 38 of the slides of Lecture 14: We would like to solve the optimization problem (OPT-4)

$$\text{Minimize}_{\theta \in \mathbb{R}^d} \quad -(\nabla_\theta L_{\theta_k}(\theta)|_{\theta=\theta_k})^T(\theta - \theta_k) \quad (2)$$

$$\text{subject to} \quad \frac{1}{2}(\theta - \theta_k)^T H(\theta - \theta_k) - \delta \leq 0. \quad (3)$$

We use θ^* to denote an optimizer of the above *primal* optimization problem (2)-(3). Note that in the above we write “Minimize” instead of “Maximize” simply to follow the conventions of the literature of optimization theory. Here we focus on the case where H is a *positive definite* matrix to avoid the technicalities (while H is only *non-negative definite* in general).

Based on the optimization theory, (OPT-4) is a convex optimization problem as both the objective and the constraints are convex functions. In this case, one standard way is to convert the constrained (OPT-4) into an unconstrained *dual* problem by defining the *Lagrangian* $\mathcal{L}(\theta, \lambda)$ and the *dual function* $D(\lambda)$ as:

$$\mathcal{L}(\theta, \lambda) := -(\nabla_\theta L_{\theta_k}(\theta)|_{\theta=\theta_k})^T(\theta - \theta_k) + \lambda \left(\frac{1}{2}(\theta - \theta_k)^T H(\theta - \theta_k) - \delta \right) \quad (4)$$

$$D(\lambda) := \min_{\theta \in \mathbb{R}^d} \mathcal{L}(\theta, \lambda), \quad (5)$$

where λ is called the *Lagrange multiplier*. Moreover, we call the following the *dual problem* of (OPT-4):

$$\max_{\lambda \geq 0} D(\lambda). \quad (6)$$

For ease of notation, we define $\lambda^* := \arg \max_{\lambda \geq 0} D(\lambda)$ as the optimizer of the dual problem. By the standard theory of convex optimization, if there exists one strictly feasible point in (OPT-4), then the optimal value of the dual problem is equal to the original problem (usually called “strong duality”). Moreover, if strong duality holds and a dual optimal solution λ^* exists, then any optimizer of the primal problem is also a minimizer of $\mathcal{L}(\theta, \lambda^*)$, i.e., $\theta^* = \arg \min_\theta \mathcal{L}(\theta, \lambda^*)$. For more details on duality, please refer to Chapter 5 of https://web.stanford.edu/~boyd/cvxbook/bv_cvxbook.pdf.

(a) In this problem, please show that the dual function $D(\lambda)$ of (OPT-4) can be written as:

$$D(\lambda) = \frac{-1}{2\lambda} ((\nabla_{\theta} L_{\theta_k}(\theta)|_{\theta=\theta_k})^T H^{-1} (\nabla_{\theta} L_{\theta_k}(\theta)|_{\theta=\theta_k})) - \lambda \delta. \quad (7)$$

Accordingly, please find out λ^* based on (7).

(b) By the λ^* found in (a) and the property $\theta^* = \arg \min_{\theta} \mathcal{L}(\theta, \lambda^*)$, show that $\theta^* = \theta_k + \alpha H^{-1} \nabla_{\theta} L_{\theta_k}(\theta)|_{\theta=\theta_k}$. What is the step size α ?

Problem 3 (Policy Gradient Algorithms With Function Approximation) (30+30=60 points)

In this problem, we will implement two policy gradient algorithms with the help of neural function approximators: (1) REINFORCE with value function as the baseline; (2) Advantage Actor-Critic algorithm (A2C). For the pseudo code of the algorithms, please see Page 19 of Lecture 8 and Page 24 of Lecture 13. You may write your code in either PyTorch or TensorFlow (though the sample code presumes PyTorch framework). If you are a beginner in learning the deep learning framework, please refer to the following tutorials:

- PyTorch: <https://pytorch.org/tutorials/>
- Tensorflow: <https://www.tensorflow.org/tutorials>

For the deliverables, please submit the following:

- Technical report: Please summarize all your experimental results in 1 single report (and please be brief)
- All your source code
- Your well-trained models (REINFORCE and A2C) saved in either .pth files or .ckpt files

(a) We start by solving the simple “CartPole-v0” problem (<https://gym.openai.com/envs/CartPole-v0/>) using the REINFORCE algorithm with value function as the baseline function. Read through `reinforce_baseline.py` and then implement the member functions of the class **Policy** and the function **train**. Please briefly summarize your results in the report and document all the hyperparameters (e.g. learning rate, NN architecture,) of your experiments.

(b) Based on the code for (a), implement the A2C algorithm by replacing the REINFORCE update with TD(0) bootstrapping and solve the “LunarLander-v2” problem, which has slightly higher state and action space dimensionality (<https://gym.openai.com/envs/LunarLander-v2/>). Save your code in another file named **a2c.py**. Please add comments to your code whenever needed for better readability. Again, please briefly summarize your results in the report and document all the hyperparameters of your experiments. (Note: As LunarLander is a more challenging environment than CartPole, it might require some efforts to tune the hyperparameters, e.g., learning rates or learning rate scheduler. You could either do grid search or even use some more advanced techniques such as the evolutionary methods. As the main purpose of this homework is to help you get familiar with RL implementation, there is NO hard requirement on the obtained returns of this LunarLander task. Just try your best and enjoy!)